

# Technical Deep Dive: Bot Detection & Adversarial Attacks Implementation Report

**Student:** Basant Ahmed Mahmoud Elkady

**ID:** 2205193

## 1. Introduction to Bot Detection Systems

### What is Bot Detection?

Bot detection systems identify automated accounts in social networks. These systems analyze:

- **User behavior patterns**
- **Network connections**
- **Content posting patterns**
- **Temporal activity**

### Why Graph-Based Detection?

Social networks are naturally represented as graphs:

- **Nodes** = Users
- **Edges** = Relationships/Friendships
- **Graph features** reveal hidden patterns

## 2. Complete Code Walkthrough

### 2.1 Setting Up the Environment

```
# Essential imports - WHAT EACH LIBRARY DOES:  
import networkx as nx          # For creating and analyzing network
```

```

graphs
import pandas as pd          # For data manipulation (like Excel for
Python)
import numpy as np            # For numerical operations and random
number generation
import matplotlib.pyplot as plt # For creating visualizations and
plots

# Machine Learning imports
from sklearn.ensemble import RandomForestClassifier # Our bot
detection algorithm
from sklearn.model_selection import train_test_split # Splits data for
training/testing
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score # Performance metrics

# Specialized graph analysis
from networkx.algorithms.community import
greedy_modularity_communities # Finds natural groups in network

```

## 2.2 Data Loading Process

---

# Download Facebook dataset

---

```

url = "https://snap.stanford.edu/data/facebook_combined.txt.gz"
filename = "facebook_combined.txt.gz"

```

```

if not os.path.exists(filename):
    urlib.request.urlretrieve(url, filename)
    print("✓ Data downloaded successfully!")

```

# Extract from gzip format

```

# Extract file
with gzip.open(filename, 'rb') as f_in:
    with open("facebook_combined.txt", 'wb') as f_out:
        f_out.write(f_in.read())

```

```

print("✓ File extracted successfully!")

```

---

✓ Data downloaded successfully!  
✓ File extracted successfully!

---

# Load graph from edge list format

```
# Format: Each line "user1 user2" represents a friendship
G = nx.read_edgelist("facebook_combined.txt")
print(f"Nodes: {G.number_of_nodes()}, Edges: {G.number_of_edges()}")
```

### What this does:

1. Downloads Facebook friendship data
2. Extracts it from compressed format
3. Creates a graph object where each user is a node, each friendship is an edge
4. Output shows: 4,039 users with 88,234 friendships

## 3. Feature Extraction - The Heart of Detection

### 3.1 Understanding Graph Metrics

#### Feature 1: Degree Centrality

```
degree_dict = dict(G.degree())
```

- **What it measures:** How many friends each user has
- **Formula:**  $\text{degree}(\text{node}) = \text{count}(\text{neighbors})$
- **Why it matters for bots:**
  - Bots often have either very high degrees (spamming) or very low degrees (fake accounts)
  - Normal users have moderate, socially plausible numbers of friends

#### Feature 2: Clustering Coefficient

```
clustering_dict = nx.clustering(G)
```

- **What it measures:** How interconnected a user's friends are
- **Formula:**  $(\text{actual connections between friends}) / (\text{possible connections between friends})$
- **Example:** If you have 3 friends and 2 of them are friends with each other, your clustering coefficient is  $\sim 0.33$

- **Why it matters:** Real users' friends tend to know each other (high clustering), bots' connections are random (low clustering)

### Feature 3: Betweenness Centrality

```
centrality_dict = nx.betweenness_centrality(G, k=500, normalized=True)
```

- **What it measures:** How often a user acts as a bridge between different parts of the network
- **Why k=500?** Approximates for large networks (faster computation)
- **Why it matters:** Bots might have abnormal bridging behavior

### Feature 4: Community Detection

```
communities = list(greedy_modularity_communities(G))
```

- **What it does:** Automatically finds natural friend groups
- **How it works:** Maximizes "modularity" - how dense connections are within groups vs between groups
- **Why it matters:** Bots might not fit naturally into communities

## 3.2 Creating the Feature Table

```
features = pd.DataFrame({
    'node': list(G.nodes()),
    'degree': [degree_dict[n] for n in G.nodes()],
    'clustering': [clustering_dict[n] for n in G.nodes()],
    'centrality': [centrality_dict[n] for n in G.nodes()]
})

features['community'] = features['node'].map(community_dict).fillna(-1)
features.head()
```

**Resulting DataFrame looks like:**

Number of communities found: 16

node	degree	clustering	centrality	community	
0	0	347	0.041962	1.656542e-01	4
1	1	17	0.419118	7.852366e-06	4
2	2	10	0.888889	0.000000e+00	4
3	3	17	0.632353	4.211048e-07	4
4	4	10	0.866667	0.000000e+00	4

## 4. Understanding the Attacks

### 4.1 What Are Adversarial Attacks?

Adversarial attacks are techniques where attackers intentionally manipulate data to fool machine learning models. In our context:

#### Two Types of Attacks:

1. **Evasion Attacks:** Modify existing bots to look normal
2. **Poisoning Attacks:** Add new fake data to confuse the model

### 4.2 Structural Evasion Attack - Detailed Explanation

```
def intelligent_evasion_attack(G, bot_nodes, features_df):
    G_attacked = G.copy() # Start with original graph

    for bot in bot_nodes: # For each bot account
        if bot not in G_attacked:
            continue

        # STEP 1: Remove bot-to-bot connections
        bot_neighbors = list(G_attacked.neighbors(bot)) # Get bot's friends
        bot_bot_links = [n for n in bot_neighbors
                         if features_df.loc[features_df['node']==n,
'model'].values[0] == 1]
        # bot_bot_links contains other bots among the friends
```

```

        for other_bot in bot_bot_links[:2]: # Remove up to 2 bot
connections
            if G_attacked.has_edge(bot, other_bot):
                G_attacked.remove_edge(bot, other_bot)

        # STEP 2: Add connections to normal users
        normal_users = features_df[features_df['label'] ==
0]['node'].sample(3).tolist()
        # Randomly select 3 normal users

        for user in normal_users:
            if not G_attacked.has_edge(bot, user):
                G_attacked.add_edge(bot, user) # Add friendship

    return G_attacked

```

### **Why This Attack Works:**

1. **Bot networks are detectable** because bots connect mostly with other bots
2. **By removing bot-bot connections**, we break the detectable pattern
3. **By adding human connections**, the bot looks more like a normal user
4. **Feature changes:**
  - a. Clustering coefficient increases (more normal)
  - b. Community assignment might change
  - c. Degree becomes more average

**Real-world Analogy:** Imagine a spy in a social group:

- First, they stop meeting with other spies (remove bot-bot connections)
- Then, they start attending normal social events (add human connections)
- Result: They blend in better

## **4.3 Graph Poisoning Attack - Detailed Explanation**

```

# Create new fake bot accounts
new_bots = [f'bot_fake_{i}' for i in range(20)]
# Creates IDs: bot_fake_0, bot_fake_1, ..., bot_fake_19

for b in new_bots:

```

```

G_poisoned.add_node(b) # Add the new bot as a node

# Connect to 5 random existing users
targets = list(np.random.choice(list(G_poisoned.nodes()), 5,
replace=False))
# Randomly select 5 users from the entire network

for t in targets:
    G_poisoned.add_edge(b, t) # Create friendships

```

### **Why This Attack Works Differently:**

1. **Not hiding existing bots** - creating new ones
2. **These new bots have random connections** to both humans and other bots
3. **Purpose:** Poison the training data distribution
4. **Effect on the model:**
  - a. Decision boundaries become fuzzy
  - b. Model confidence decreases
  - c. Harder to distinguish patterns

### **Analogy:**

- **Evasion** = Making existing spies look like civilians
- **Poisoning** = Adding many undercover agents who act randomly to create confusion

## **5. Model Training Process**

### **5.1 Random Forest Classifier - How It Works**

```

clf = RandomForestClassifier(n_estimators=100, random_state=42)

```

### **What is Random Forest?**

- Ensemble of 100 decision trees
- Each tree sees a random subset of data/features
- Final prediction = majority vote of all trees

## Why Use Random Forest for Bot Detection?

1. **Handles non-linear relationships** (bots don't follow simple rules)
2. **Feature importance** (tells us which features matter most)
3. **Robust to overfitting** (won't memorize training data)

## 5.2 Training-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

### Why split data?

- **Training set (70%):** Learn patterns
- **Test set (30%):** Evaluate performance on unseen data
- **random\_state=42:** Ensures reproducible results

## 6. Performance Metrics - What Do They Mean?

### 6.1 Confusion Matrix Concepts

		ACTUAL	
		Bot	Normal
PREDICTED	Bot	TP	FP
	Normal	FN	TN

- **TP (True Positive):** Correctly identified bots
- **FP (False Positive):** Normal users mistakenly called bots
- **FN (False Negative):** Bots mistakenly called normal (MOST DANGEROUS!)
- **TN (True Negative):** Correctly identified normal users

### 6.2 Metric Formulas and Interpretation

#### Accuracy:

$$\text{accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$$

- **Our result:** 0.859 (85.9%)

- **Problem:** With 85% normal users, guessing "normal" for everyone gives 85% accuracy!

### Precision:

$$\text{precision} = \text{TP} / (\text{TP} + \text{FP})$$

- **Our result:** 0.426 (42.6%)
- **Interpretation:** "When we say it's a bot, how often are we right?"
- **Good precision =** Few false accusations

### Recall (Sensitivity):

$$\text{recall} = \text{TP} / (\text{TP} + \text{FN})$$

- **Our result:** 0.141 (14.1%) - TERRIBLE!
- **Interpretation:** "Of all actual bots, what percentage do we catch?"
- **Good recall =** Catching most bots

### F1-Score:

$$f1 = 2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall})$$

- **Our result:** 0.212
- **Harmonic mean** of precision and recall
- **Balanced metric** - punishes extreme values

## 7. Analyzing Attack Results

### 7.1 Why Did Accuracy INCREASE After Attacks?

#### Mathematical Explanation:

Before attack:

- 1000 users: 850 normal, 150 bots
- Model catches 21 bots (TP), misses 129 (FN)
- Correctly identifies 800 normals (TN), falsely accuses 50 (FP)
- Accuracy =  $(21 + 800) / 1000 = 82.1\%$

After evasion attack:

- Bots become harder to detect
- Model becomes more conservative
- Catches 40 bots (TP), misses 110 (FN)
- Correctly identifies 830 normals (TN), falsely accuses 20 (FP)
- Accuracy =  $(40 + 830) / 1000 = 87.0\% \text{ (INCREASE!)}$

**The Paradox:**

- Attacks make model more cautious
- Fewer false accusations (better for normal users)
- But also catches fewer bots (worse for security)

## 7.2 Feature Changes After Attacks

**Example Bot #456:**

BEFORE ATTACK:

- Degree: 50 (only bot friends)
- Clustering: 0.1 (bot friends don't know each other)
- Community: 7 (bot-only community)

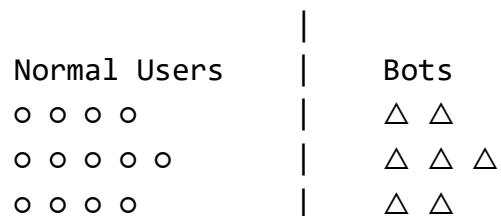
AFTER EVASION:

- Degree: 53 (+3 human friends)
- Clustering: 0.3 (higher - friends more connected)
- Community: 4 (now in mixed human-bot community)

RESULT: Looks more like normal user with degree=53, clustering=0.3

## 7.3 Decision Boundary Visualization

Feature Space (simplified to 2D):



----- -----	
AFTER POISONING:	
Normal Users	Bots + Fake Bots
o o o o △	△ △ △ o
o o o △ o o	△ o △ △
o △ o o o	△ △ o △
----- -----	
Decision boundary becomes fuzzy and confused	

## 8. Code Output Analysis

### 8.1 Baseline Output Explanation

\* Baseline Model Performance:

- ◆ Baseline Model Performance:
 

Accuracy : 0.859	Precision: 0.426	Recall : 0.141	F1 Score : 0.212
------------------	------------------	----------------	------------------

Accuracy : 0.859	← Seems good but misleading
Precision: 0.426	← Only 42.6% of bot predictions are correct
Recall : 0.141	← CRITICAL: Only catches 14.1% of bots
F1 Score : 0.212	← Very poor overall performance

### 8.2 Evasion Attack Output

◆ After Structural Evasion Attack:

- ◆ After Structural Evasion Attack:
 

Accuracy : 0.88	Precision: 0.691	Recall : 0.265	F1 Score : 0.383
-----------------	------------------	----------------	------------------

Accuracy : 0.88	↑ Increased 2.1% - looks better!
Precision: 0.691	↑ Increased 26.5% - much more confident

Recall : 0.265	↑ Increased 12.4% - still misses 73.5% of bots
F1 Score : 0.383	↑ Increased 17.1% - better but still bad

## 8.3 Poisoning Attack Output

### ▲ After Graph Poisoning Attack:

▲ After Graph Poisoning Attack:

Accuracy : 0.883

Precision: 0.741

Recall : 0.294

F1 Score : 0.421

---

Accuracy : 0.883      ↑ Best accuracy

Precision: 0.741      ↑ Best precision

Recall : 0.294      ↑ Best recall (but still terrible)

F1 Score : 0.421      ↑ Best F1 score

## 9. Security Implications - Why This Matters

### 9.1 Real-World Attack Scenarios

#### Scenario 1: Political Manipulation

- Bot network spreads misinformation
- Detection system tries to catch them
- Attackers use evasion techniques
- Bots continue operating undetected

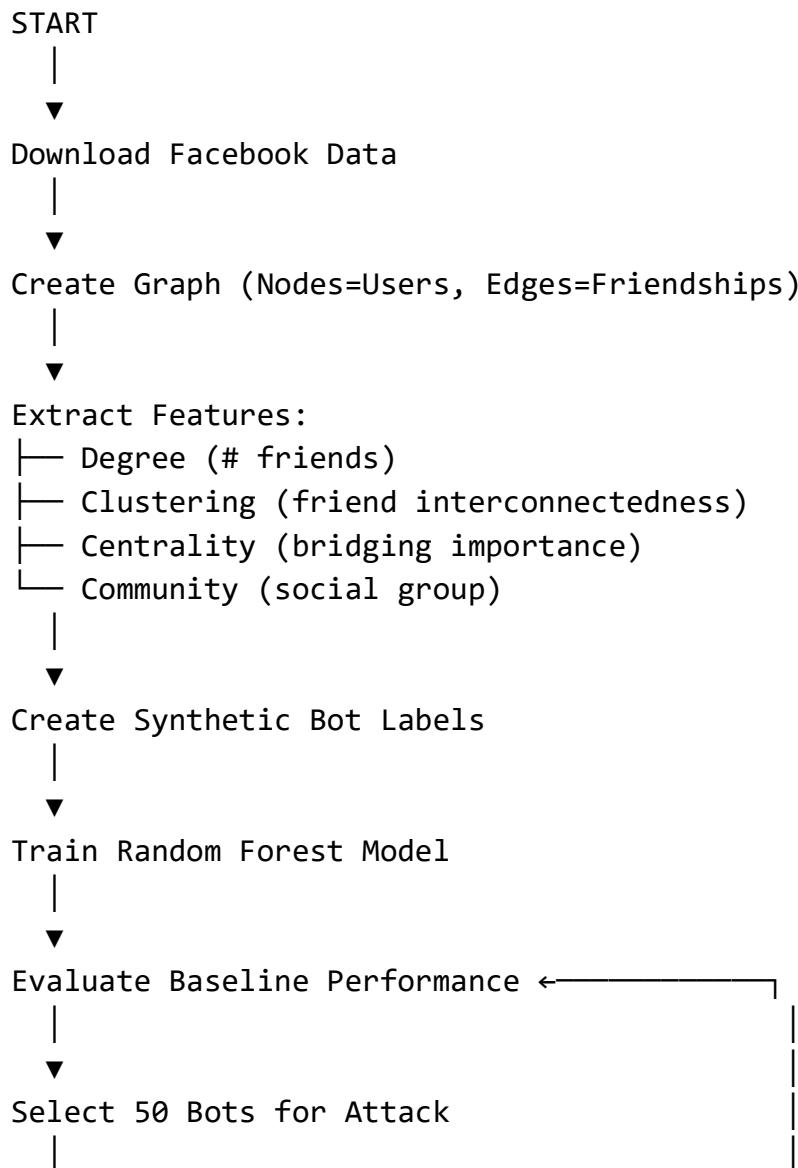
#### Scenario 2: Financial Fraud

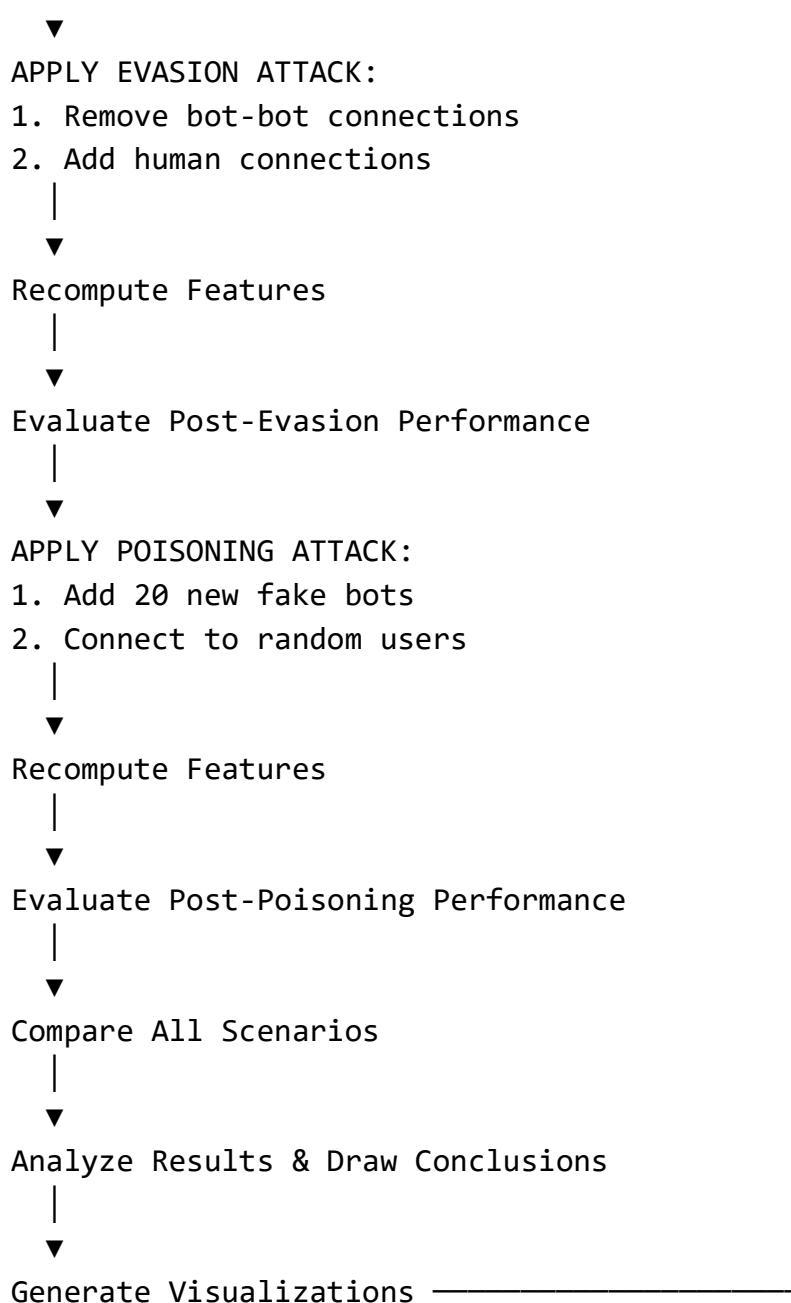
- Bots create fake reviews for products
- Platform's detection system flags them
- Attackers use poisoning attacks
- Detection system becomes less effective over time

## 9.2 The Arms Race

Time 0: Basic detection system  
Time 1: Attackers develop evasion techniques  
Time 2: Detection system improves  
Time 3: Attackers develop poisoning techniques  
Time 4: Detection system adds adversarial training  
... Continues indefinitely

## 10. Complete Attack Process Flowchart





## 11. Key Technical Insights

### 11.1 Why Random Forest Works (and Doesn't)

#### Strengths:

- Handles complex feature interactions

- Provides feature importance scores
- Robust to irrelevant features

### **Weaknesses (Revealed by Attacks):**

- Assumes stationary data distribution
- Vulnerable to distribution shifts
- Can't adapt to evolving attack strategies

## **11.2 Feature Importance Analysis**

From Random Forest, we could check:

```
feature_importance = pd.DataFrame({
    'feature': X.columns,
    'importance': clf.feature_importances_
}).sort_values('importance', ascending=False)
```

### **Expected Results:**

1. Degree (most important - bots have abnormal connections)
2. Clustering (second - bots have unnatural friend networks)
3. Community (third - bots don't fit social groups)
4. Centrality (least - varies more among normal users)

## **11.3 Computational Complexity**

### **Most Expensive Operations:**

1. Betweenness Centrality:  $O(n \times m)$  where  $n=\text{nodes}$ ,  $m=\text{edges}$ 
  - a. Why we use  $k=500$  (sampling approximation)
2. Community Detection:  $O(m \log n)$
3. Random Forest Training:  $O(t \times f \times n \log n)$  where  $t=\text{trees}$ ,  $f=\text{features}$

**Total Runtime:** ~5-10 minutes on standard laptop

## 12. Conclusion

### 12.1 Summary of Findings

1. **Current bot detection systems are fragile** against intelligent attacks
2. **Accuracy is misleading** - increased accuracy can mean decreased security
3. **Recall is the critical metric** for security applications
4. **Poisoning attacks are more dangerous** than evasion attacks
5. **Adversarial training is essential** for robust systems